

PREDICTIVE ANALYSIS IN PROPERTY VALUATION.

The Outliers Team Members

Jonquil Phelan

Brenda Ngigi

George Mbugua

Vincent Kiplangat

Kenneth Gakuru

Fanice Andayi

Wallace Ouma

Ian Odhiambo

Charles Gaithuma

A)INTRODUCTION

In the dynamic landscape of real estate, the ability to provide clients with accurate and data-driven insights is crucial for success. Real estate agencies play a pivotal role in guiding homeowners through key decisions such as buying, selling, or renting properties. Some of the key decisions that real estate agencies assist homeowners with include: pricing strategies, Market trends and analysis and property inspections. The aim of this project is to empower real estate agencies with a powerful tool – a regression-based model – that can predict the potential increase in property value based on various features influenced by property characteristics such as number of bedrooms, year built, Number of floors in the home, Total living space area in square feet, Overall condition of the home and the location of the home. By leveraging this model, agencies can offer tailored advice to their clients, ensuring informed decisions that maximize the return on investment in the competitive real estate market.

B)BACKGROUND OF THE DATA

The housing market in King County, Washington, has a rich history and has experienced significant growth and changes over the years. King County is situated in the northwestern part of the United States and encompasses the city of Seattle, which is a major economic and cultural hub in the region. As the region's economy thrived, King County experienced substantial population growth. This influx of residents led to increased demand for housing, both in urban and suburban areas. Seattle, with its iconic skyline, became a sought-after destination for tech professionals and urban enthusiasts. The real estate market in King County is known for its competitiveness. The region's diverse neighborhoods offer a range of housing options, from historic homes in older districts to modern developments in suburban areas.

C) PROBLEM STATEMENT

What is the prevailing circumstance? In identifying the prevailing circumstance, real estate agencies wrestle with the ongoing challenge of delivering insightful guidance to clients regarding pricing strategies, Market trends and analysis and property inspections. Clients frequently seek advice on buying, selling, or renting properties. The complexity arises from the multitude of factors influencing property value, making it challenging to accurately quantify the best price for a particular property.

What problem are we trying to solve? The problem at hand revolves around the need for a comprehensive and precise solution to guide real estate agencies in offering informed recommendations to their clients. To tackle this challenge, this project is centered on the development of a regression-based approach, leveraging the King County House Sales dataset. Through this approach, the objective is to construct a robust predictive model that takes into account various features, including property size, condition, location, and more. This model is designed to empower real estate agencies in providing evidence-based insights, pinpointing the property features most likely to yield a substantial return on investment.

How the project aims to solve the problem? In essence, this project aims to equip real estate agencies with a valuable tool that goes beyond enhancing advisory capabilities—it positions them as trusted partners in their clients' real estate journeys. By developing and implementing the regression model, this project strives to optimize property value and contribute to the overall success of real estate agencies, solidifying their client-centric approach in the dynamic real estate landscape.

D) OBJECTIVES

Main Objective:

The main objective of this project is to develop a predictive regression model that assists real estate agencies in advising clients on house prices. The model aims to predict the potential variation in property value based on property characteristics, providing valuable insights to guide clients in making informed decisions about their investments.

Specific Objectives

- i). Identify Key Factors Influencing House Prices in King County, California, to provide valuable insights for precise pricing strategies.
- ii). Analyze Model Performance using metrics such as mean squared error, R-squared values, and residual analysis to gauge the model's effectiveness.
- iii). Provide Actionable Recommendations to the Real Estate Agency for improving profitability and market presence, leveraging insights from the model.

E) NOTEBOOK STRUCTURE

1. Overview

2. Business Understanding

- 3.Data Understanding
- 4.Data Cleaning
- 5.Statistical Analysis
- 6.Data Preparation
- 7.Modelling
- 8.Regression Results
- 9.Conclusion - Recommendations,Limitations and Next Steps

1.OVERVIEW

This project is centered on the development of a multiple linear regression model with the primary objective of predicting property prices in the real estate market. The focal point of this analysis is the dependent variable, "price," while the independent variables encompass a comprehensive range of property characteristics. These characteristics include, but are not limited to, the number of bedrooms, year built, number of floors in the home, total living space area in square feet, overall condition of the home, and the geographical location of the property.

2.BUSINESS UNDERSTANDING

Stakeholders and Their Interests:

The key stakeholders impacted by this project are Real Estate Agencies. Their interest is accurate pricing and having a competitive advantage and the value they seek is Improved decision making and client satisfaction.

Success Criteria:

Establish measurable success criteria that align with the business objectives. For instance, success might be defined by the model's accuracy in predicting property prices within a certain margin.

```
# Importing necessary libraries for data analysis and visualization

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt # for data visualization.
from pandas.api.types import is_numeric_dtype # Used to check if a
data type is numeric.
%matplotlib inline
import seaborn as sns # for enhanced data visualization.

from statsmodels.stats.outliers_influence import
variance_inflation_factor # For calculating Variance Inflation Factor
(VIF).
from statsmodels.graphics.regressionplots import plot_partregress_grid
```

```

# For partial regression plots.
from sklearn.model_selection import train_test_split # Used to split
data into training and testing sets.
from sklearn.feature_selection import RFE # Recursive Feature
Elimination for feature selection.
from sklearn.preprocessing import StandardScaler #
Standardizing/Scaling features.
from sklearn.preprocessing import PolynomialFeatures # Generate
polynomial features.
from sklearn.linear_model import LinearRegression # Linear Regression
model.
from sklearn.metrics import mean_squared_error, r2_score # Evaluation
metrics for model performance.
import statsmodels.api as sm
from scipy.stats import kstest

# Statsmodels is used to create statistical models.
from scipy import stats # Scientific computing library for
statistical tests.
from scipy.stats import f_oneway # One-way ANOVA statistical test.
from scipy.stats import ttest_ind # Independent sample t-test for
comparing means.
import warnings # handle warnings during code execution.
warnings.filterwarnings("ignore") # Ignore warnings to improve code
readability.

#loading the dataset
Data= 'kc_house_data.csv'
# Load the CSV file into a Pandas DataFrame
kc_house_data = pd.read_csv(Data)
# Display the first few rows of the dataset to ensure it's loaded
correctly
kc_house_data.head()

```

	id	date	price	bedrooms	bathrooms	sqft_living
0	7129300520	10/13/2014	221900.0	3	1.00	1180
1	6414100192	12/9/2014	538000.0	3	2.25	2570
2	5631500400	2/25/2015	180000.0	2	1.00	770
3	2487200875	12/9/2014	604000.0	4	3.00	1960
4	1954400510	2/18/2015	510000.0	3	2.00	1680

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
0	5650	1.0	NaN	NONE	...	7 Average	1180	
1	7242	2.0	NO	NONE	...	7 Average	2170	

2	10000	1.0	NO	NONE	...	6	Low	Average	770
3	5000	1.0	NO	NONE	...	7	Average	1050	
4	8080	1.0	NO	NONE	...	8	Good	1680	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
0	0.0	1955	0.0	98178	47.5112	-122.257	
1	400.0	1951	1991.0	98125	47.7210	-122.319	
2	0.0	1933	NaN	98028	47.7379	-122.233	
3	910.0	1965	0.0	98136	47.5208	-122.393	
4	0.0	1987	0.0	98074	47.6168	-122.045	

	sqft_living15	sqft_lot15
0	1340	5650
1	1690	7639
2	2720	8062
3	1360	5000
4	1800	7503

[5 rows x 21 columns]

3.DATA UNDERSTANDING

-The data utilized for this project has been sourced from Kaggle

Data Source:

The dataset used for this project is the King County House Sales dataset, which is available in the kc_house_data.csv file. The dataset contains information about house sales in King County, providing details such as property features, location, sale prices, and renovation-related variables.

Data Size:

The kc_house_dataset contains 21597 rows and 21 columns.The columns include :

Column Names and Descriptions for King County Data Set

- **id** - Unique identifier for a house
- **date** - Date house was sold
- **price** - Sale price (prediction target)
- **bedrooms** - Number of bedrooms
- **bathrooms** - Number of bathrooms
- **sqft_living** - Square footage of living space in the home

- `sqft_lot` - Square footage of the lot
- `floors` - Number of floors in house
- `waterfront` - Whether the house is on a waterfront
 - Includes Duwamish, Elliott Bay, Puget Sound, Lake Union, Ship Canal, Lake Washington, Lake Sammamish, other lake, and river/slough waterfronts
- `view` - Quality of view from house
 - Includes views of Mt. Rainier, Olympics, Cascades, Territorial, Seattle Skyline, Puget Sound, Lake Washington, Lake Sammamish, small lake / river / creek, and other
- `condition` - How good the overall condition of the house is. Related to maintenance of house.
- `grade` - Overall grade of the house. Related to the construction and design of the house.
- `sqft_above` - Square footage of house apart from basement
- `sqft_basement` - Square footage of the basement
- `yr_built` - Year when house was built
- `yr_renovated` - Year when house was renovated
- `zipcode` - ZIP Code used by the United States Postal Service
- `lat` - Latitude coordinate
- `long` - Longitude coordinate
- `sqft_living15` - The square footage of interior housing living space for the nearest 15 neighbors
- `sqft_lot15` - The square footage of the land lots of the nearest 15 neighbors

Data Types:

The data types include categorical and numerical variables. Columns with categorical variables include `date`, `waterfront`, `view`, `condition`, `grade` and `sqft_basement` while columns with numerical variables include

`id`, `price`, `bedrooms`, `bathrooms`, `sqft_living`, `sqft_lot`, `floors`, `yr_built`, `yr_renovated`, `zipcode`, `lat`, `long`, `sqft_above`, `sqft_living15` and `sqft_lot15`.

a) Determining the number of records

```
num_records = kc_house_data.shape
print("Number of records:", num_records)
```

Number of records: (21597, 21)

- The data contains 21597 rows and 21 columns

b) Preview top and bottom of our dataset

```
# Preview the top of the dataset
```

```
top_rows = kc_house_data.head()
```

```
top_rows
```

\	id	date	price	bedrooms	bathrooms	sqft_living
0	7129300520	10/13/2014	221900.0	3	1.00	1180
1	6414100192	12/9/2014	538000.0	3	2.25	2570
2	5631500400	2/25/2015	180000.0	2	1.00	770
3	2487200875	12/9/2014	604000.0	4	3.00	1960
4	1954400510	2/18/2015	510000.0	3	2.00	1680

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
0	5650	1.0	NaN	NONE	...	7 Average	1180	
1	7242	2.0	NO	NONE	...	7 Average	2170	
2	10000	1.0	NO	NONE	...	6 Low Average	770	
3	5000	1.0	NO	NONE	...	7 Average	1050	
4	8080	1.0	NO	NONE	...	8 Good	1680	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
0	0.0	1955	0.0	98178	47.5112	-122.257	
1	400.0	1951	1991.0	98125	47.7210	-122.319	
2	0.0	1933	NaN	98028	47.7379	-122.233	
3	910.0	1965	0.0	98136	47.5208	-122.393	
4	0.0	1987	0.0	98074	47.6168	-122.045	

	sqft_living15	sqft_lot15
0	1340	5650
1	1690	7639
2	2720	8062
3	1360	5000
4	1800	7503

```
[5 rows x 21 columns]
```

```
# Preview the bottom of the dataset
```

```
bottom_rows = kc_house_data.tail()
```

```
bottom_rows
```

	id	date	price	bedrooms	bathrooms
sqft_living \					
21592	263000018	5/21/2014	360000.0	3	2.50
1530					
21593	6600060120	2/23/2015	400000.0	4	2.50
2310					
21594	1523300141	6/23/2014	402101.0	2	0.75
1020					
21595	291310100	1/16/2015	400000.0	3	2.50
1600					
21596	1523300157	10/15/2014	325000.0	2	0.75
1020					

	sqft_lot	floors	waterfront	view	...	grade	sqft_above \
21592	1131	3.0	NO	NONE	...	8 Good	1530
21593	5813	2.0	NO	NONE	...	8 Good	2310
21594	1350	2.0	NO	NONE	...	7 Average	1020
21595	2388	2.0	NaN	NONE	...	8 Good	1600
21596	1076	2.0	NO	NONE	...	7 Average	1020

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long
\						
21592	0.0	2009	0.0	98103	47.6993	-122.346
21593	0.0	2014	0.0	98146	47.5107	-122.362
21594	0.0	2009	0.0	98144	47.5944	-122.299
21595	0.0	2004	0.0	98027	47.5345	-122.069
21596	0.0	2008	0.0	98144	47.5941	-122.299

	sqft_living15	sqft_lot15
21592	1530	1509
21593	1830	7200
21594	1020	2007
21595	1410	1287
21596	1020	1357

[5 rows x 21 columns]

c).Checking data types in various columns

- This involves checking whether the columns have appropriate data types

```
data_types = kc_house_data.dtypes
print(f>Data types of each column:\n{data_types})
```


Data types of each column:

```
id          int64
date        object
price       float64
bedrooms    int64
bathrooms   float64
sqft_living int64
sqft_lot    int64
floors       float64
waterfront  object
view        object
condition   object
grade       object
sqft_above  int64
sqft_basement object
yr_built    int64
yr_renovated float64
zipcode     int64
lat         float64
long        float64
sqft_living15 int64
sqft_lot15  int64
dtype: object
```

-The columns have three data types:

- Integers which include id,bedrooms,sqft_living,sqft_lot,sqft_above,yr built,zipcode,sqft living15,sqft lot15
- Float data types include price,bathrooms,floors,year renovated,latitudes and longitudes
- Object data type include the columns date,waterfront,view,condition,grade and sqft basement.

d)Descriptive statistics

```
kc_house_data.describe()
```

	id	price	bedrooms	bathrooms
count	2.159700e+04	2.159700e+04	21597.000000	21597.000000
mean	4.580474e+09	5.402966e+05	3.373200	2.115826
std	2.876736e+09	3.673681e+05	0.926299	0.768984
min	1.000102e+06	7.800000e+04	1.000000	0.500000
25%	2.123049e+09	3.220000e+05	3.000000	1.750000

```

1430.000000
50%    3.904930e+09  4.500000e+05      3.000000      2.250000
1910.000000
75%    7.308900e+09  6.450000e+05      4.000000      2.500000
2550.000000
max     9.900000e+09  7.700000e+06     33.000000      8.000000
13540.000000

```

```

          sqft_lot      floors      sqft_above      yr_built
yr_renovated \
count  2.159700e+04  21597.000000  21597.000000  21597.000000
17755.000000
mean   1.509941e+04      1.494096   1788.596842   1970.999676
83.636778
std    4.141264e+04      0.539683    827.759761    29.375234
399.946414
min     5.200000e+02      1.000000    370.000000   1900.000000
0.000000
25%     5.040000e+03      1.000000   1190.000000   1951.000000
0.000000
50%     7.618000e+03      1.500000   1560.000000   1975.000000
0.000000
75%     1.068500e+04      2.000000   2210.000000   1997.000000
0.000000
max     1.651359e+06      3.500000   9410.000000  2015.000000
2015.000000

```

```

          zipcode          lat          long  sqft_living15
sqft_lot15
count  21597.000000  21597.000000  21597.000000  21597.000000
21597.000000
mean   98077.951845    47.560093   -122.213982   1986.620318
12758.283512
std     53.513072      0.138552      0.140724    685.230472
27274.441950
min     98001.000000    47.155900   -122.519000    399.000000
651.000000
25%     98033.000000    47.471100   -122.328000   1490.000000
5100.000000
50%     98065.000000    47.571800   -122.231000   1840.000000
7620.000000
75%     98118.000000    47.678000   -122.125000   2360.000000
10083.000000
max     98199.000000    47.777600   -121.315000   6210.000000
871200.000000

```

e)Summary of our dataframe

```
kc_house_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21597 non-null  int64
1   date                 21597 non-null  object
2   price               21597 non-null  float64
3   bedrooms            21597 non-null  int64
4   bathrooms           21597 non-null  float64
5   sqft_living         21597 non-null  int64
6   sqft_lot            21597 non-null  int64
7   floors              21597 non-null  float64
8   waterfront          19221 non-null  object
9   view                21534 non-null  object
10  condition            21597 non-null  object
11  grade               21597 non-null  object
12  sqft_above          21597 non-null  int64
13  sqft_basement       21597 non-null  object
14  yr_built            21597 non-null  int64
15  yr_renovated        17755 non-null  float64
16  zipcode             21597 non-null  int64
17  lat                 21597 non-null  float64
18  long                21597 non-null  float64
19  sqft_living15       21597 non-null  int64
20  sqft_lot15          21597 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB

```

4.DATA PRE-PROCESSING

Data cleaning and preparation involves addressing issues related to the quality of the dataset. It aims to ensure that the data is accurate, consistent, and free from errors. Here are some data cleaning and preparation methods we engaged in:

Handling Missing Values:

Missing values were identified and addressed depending on the type of column whether categorical or numerical. Options included removal of rows or columns with missing values, or treating missing values as a separate category.

Handling Duplicates:

Any duplicate records in the dataset were identified and addressed to avoid redundancy and potential bias in the analysis especially in the id column since it's a unique identifier.

Dealing with placeholders:

Investigation and rectification of any placeholders in the data that may affect the accuracy of the model was done.

Transforming data: Feature engineering and ordinal encoding on the data was performed.

Handling Outliers: Outliers could skew statistical analysis. Methods like trimming have been employed to address extreme values that could distort insights into housing prices attributes.

a) Handling Missing Values

The columns waterfront, view and the year renovated contain missing values which need to be addressed.

```
#checking for missing values
kc_house_data.isnull().sum()

id                0
date              0
price             0
bedrooms          0
bathrooms         0
sqft_living       0
sqft_lot          0
floors            0
waterfront       2376
view              63
condition         0
grade             0
sqft_above        0
sqft_basement     0
yr_built          0
yr_renovated      3842
zipcode           0
lat               0
long              0
sqft_living15     0
sqft_lot15        0
dtype: int64
```

i) Handling missing values in categorical columns

The categorical columns are waterfront which contains 2376 missing values and view which contains 63 missing values .

```
# Create a new dataframe of the raw data to clean
Cleaned_Data = kc_house_data.copy()

# Change waterfront missing value to NO, then to binary values.
Cleaned_Data.loc[kc_house_data.waterfront.isnull(), 'waterfront'] =
'NO'
Cleaned_Data['waterfront'] = Cleaned_Data['waterfront'].apply(lambda
x: 0 if x == 'NO' else 1)
```

```
# Change view missing value to NONE, then to numerical ordered values.
Cleaned_Data.loc[kc_house_data.view.isnull(), 'view'] = "NONE"
view_dict = {'NONE':0, 'FAIR':1, 'AVERAGE':2, 'GOOD':3, 'EXCELLENT':4}
Cleaned_Data['view'].replace(view_dict, inplace=True)
```

JUSTIFICATION

Handling 'waterfront' missing values:

Replace missing values in 'waterfront' with 'NO' since the missing values indicate that the property does not have a waterfront. Applied a binary encoding, converting 'NO' to 0 and other values (potentially 'YES') to 1 since the presence or absence of a waterfront is a significant predictor.

Handling 'view' Missing Values:

Replace missing values in 'view' with 'NONE' since the missing values imply that the property has no specific view. Replace categorical values in 'view' with numerical values .It helps capture the ordinal nature of the 'view' categories in predictive models.

ii)Handling missing values in numerical columns

The numerical column is the year renovated with 3842 missing values .

```
# Change yr_renovated missing values to 0 and add renovated column
Cleaned_Data.loc[kc_house_data.yr_renovated.isnull(), 'yr_renovated']
= 0
Cleaned_Data['renovated'] = Cleaned_Data['yr_renovated'].apply(lambda
x: 0 if x == 0 else 1)
```

JUSTIFICATION

Handling 'yr_renovated' Missing Values:

Impute missing values in 'yr_renovated' with 0 since the missing value indicates that the property has not been renovated. Create a new binary column 'renovated' based on the values in the 'yr_renovated' column. If 'yr_renovated' is 0, the 'renovated' column is set to 0; otherwise, it is set to 1.

b)Dealing with duplicates

The id column which is a unique identifier for a house was checked for duplicates.

```
# checking for duplicates using the id column
duplicates = kc_house_data[kc_house_data.duplicated(
subset = "id")]
duplicates.head()
```

	id	date	price	bedrooms	bathrooms
sqft_living \					
94	6021501535	12/23/2014	700000.0	3	1.50

```

1580
314 4139480200 12/9/2014 1400000.0 4 3.25
4290
325 7520000520 3/11/2015 240500.0 2 1.00
1240
346 3969300030 12/29/2014 239900.0 4 1.00
1000
372 2231500030 3/24/2015 530000.0 4 2.25
2180

sqft_lot floors waterfront view ... grade sqft_above
\
94 5000 1.0 NO NONE ... 8 Good 1290
314 12103 1.0 NO GOOD ... 11 Excellent 2690
325 12092 1.0 NO NONE ... 6 Low Average 960
346 7134 1.0 NO NONE ... 6 Low Average 1000
372 10754 1.0 NO NONE ... 7 Average 1100

sqft_basement yr_built yr_renovated zipcode lat
long \
94 290.0 1939 0.0 98117 47.6870 -122.386
314 1600.0 1997 0.0 98006 47.5503 -122.102
325 280.0 1922 1984.0 98146 47.4957 -122.352
346 0.0 1943 NaN 98178 47.4897 -122.240
372 1080.0 1954 0.0 98133 47.7711 -122.341

sqft_living15 sqft_lot15
94 1570 4500
314 3860 11244
325 1820 7460
346 1020 7138
372 1810 6929

[5 rows x 21 columns]

duplicates.shape

(177, 21)

# dropping the duplicated data
Cleaned_Data.drop_duplicates(subset="id", keep="first", inplace=True)

```

JUSTIFICATION

The 'id' column is a unique identifier for each property. Duplicate entries can introduce inconsistencies and bias in data analysis and modeling. Removing duplicates leads to a more accurate representation of the dataset.

c) Dealing with placeholders

The sqft_basement column contained a large number of placeholders which had to be addressed.

```
# Add has_basement column that is a binary value.
Cleaned_Data['sqft_basement'] =
Cleaned_Data['sqft_basement'].replace('?', '0').astype('float')
Cleaned_Data['basement'] = Cleaned_Data['sqft_basement'].apply(lambda
x: 0 if x == 0 else 1)
```

JUSTIFICATION

The use of '?' as a placeholder in the 'sqft_basement' column indicates missing or unknown values. The imputation and binary encoding provide consistency in handling missing or placeholder values. The resulting 'basement' column enhances the interpretability of the dataset.

d) Transforming data

```
# Change to datetime and add month column
Cleaned_Data['date'] = pd.to_datetime(Cleaned_Data['date'])
Cleaned_Data['month'] = pd.DatetimeIndex(Cleaned_Data['date']).month

# Change condition to numerical ordered values.
cond_dict = {'Poor':0, 'Fair':1, 'Average':2, 'Good':3, 'Very Good':4}
Cleaned_Data['condition'].replace(cond_dict, inplace=True)

# Change grade to numerical ordered values.
Cleaned_Data['grade'] = Cleaned_Data['grade'].map(lambda x:
int(x.split(' ')[0]))

# Add house_age column
Cleaned_Data['age'] = Cleaned_Data['date'].dt.year -
Cleaned_Data['yr_built']
```

Cleaned_Data

	id	date	price	bedrooms	bathrooms
sqft_living \					
0	7129300520	2014-10-13	221900.0	3	1.00
1180					
1	6414100192	2014-12-09	538000.0	3	2.25
2570					

2770	5631500400	2015-02-25	180000.0	2	1.00		
31960	2487200875	2014-12-09	604000.0	4	3.00		
41680	1954400510	2015-02-18	510000.0	3	2.00		
...	
215921530	2630000018	2014-05-21	360000.0	3	2.50		
215932310	6600060120	2015-02-23	400000.0	4	2.50		
215941020	1523300141	2014-06-23	402101.0	2	0.75		
215951600	291310100	2015-01-16	400000.0	3	2.50		
215961020	1523300157	2014-10-15	325000.0	2	0.75		
	sqft_lot	floors	waterfront	view	...	yr_renovated	zipcode
0	5650	1.0	0	0	...	0.0	98178
1	7242	2.0	0	0	...	1991.0	98125
2	10000	1.0	0	0	...	0.0	98028
3	5000	1.0	0	0	...	0.0	98136
4	8080	1.0	0	0	...	0.0	98074
...
21592	1131	3.0	0	0	...	0.0	98103
21593	5813	2.0	0	0	...	0.0	98146
21594	1350	2.0	0	0	...	0.0	98144
21595	2388	2.0	0	0	...	0.0	98027
21596	1076	2.0	0	0	...	0.0	98144
	lat	long	sqft_living15	sqft_lot15	renovated		
basement \ 00	47.5112	-122.257	1340	5650	0		
11	47.7210	-122.319	1690	7639	1		

2	47.7379	-122.233	2720	8062	0
0					
3	47.5208	-122.393	1360	5000	0
1					
4	47.6168	-122.045	1800	7503	0
0					
...
.					
21592	47.6993	-122.346	1530	1509	0
0					
21593	47.5107	-122.362	1830	7200	0
0					
21594	47.5944	-122.299	1020	2007	0
0					
21595	47.5345	-122.069	1410	1287	0
0					
21596	47.5941	-122.299	1020	1357	0
0					
	month	age			
0	10	59			
1	12	63			
2	2	82			
3	12	49			
4	2	28			
...			
21592	5	5			
21593	2	1			
21594	6	5			
21595	1	11			
21596	10	6			
[21420 rows x 25 columns]					

JUSTIFICATION

1.Handling the 'date' column

Convert the 'date' column to datetime format ensures that the data type is consistent and allows for convenient handling of date-related operations.

2.Handling the 'condition' column

Replace categorical values in the 'condition' column with numerical ordered values facilitates a consistent representation and captures the ordinal nature of condition ratings.

3.Handling 'grade' Column:


Parsing the 'grade' column to extract the numerical part and converting it to an integer ensures a consistent numerical representation.

4.Adding 'house_age' Column:

Creating a new 'age' column by subtracting the year built from the year of sale provides valuable information about the age of each property at the time of sale.

e)Handling outliers


```
#PRICE  
# Let's look at the price  
sns.boxplot(Cleaned_Data['price']);
```



```
# Filtering out outliers of price in the data  
clean_data = Cleaned_Data[(Cleaned_Data.price <  
Cleaned_Data.price.quantile(.995))  
                           & (Cleaned_Data.price >  
Cleaned_Data.price.quantile(.005))]
```

```
#BEDROOM  
# Let's look at the bedrooms  
sns.boxplot(Cleaned_Data['bedrooms'])
```

```
<Axes: ylabel='bedrooms'>
```



```
# the outlier in bedrooms column change to 3 bedrooms, likely due to a
typographic error
Cleaned_Data.loc[(Cleaned_Data.bedrooms == 33), 'bedrooms'] = 3
```

JUSTIFICATION

Outliers can influence the parameters of statistical models, leading to biased results. Addressing outliers present in the model for the 'price' column filtering them out on the basis of the interquartile range which helps prevent the model from being overly influenced by extreme values, leading to more accurate and robust predictions. For the outlier in bedrooms column change to 3 bedrooms, likely due to a typographic error.

5. STATISTICAL ANALYSIS

Statistical analysis is a critical component of understanding relationships within the dataset, identifying patterns, and gaining insights. In the context of a regression modeling project for predicting property value based on home renovations, here are the key steps in statistical analysis:

a) Descriptive Statistics:

- Understanding the basic characteristics of the dataset.
- **Tasks:** -Summary statistics such as mean, median, standard deviation, and quartiles for numeric variables were computed.

```
Cleaned_Data.describe()
```

	id		date	price \
count	2.142000e+04		21420	2.142000e+04
mean	4.580940e+09	2014-10-28 05:03:51.932773120		5.407393e+05
min	1.000102e+06	2014-05-02 00:00:00		7.800000e+04
25%	2.123537e+09	2014-07-21 00:00:00		3.225000e+05
50%	3.904921e+09	2014-10-15 00:00:00		4.500000e+05
75%	7.308900e+09	2015-02-13 00:00:00		6.450000e+05
max	9.900000e+09	2015-05-27 00:00:00		7.700000e+06
std	2.876761e+09		NaN	3.679311e+05

	bedrooms	bathrooms	sqft_living	sqft_lot
floors \				
count	21420.000000	21420.000000	21420.000000	2.142000e+04
21420.000000				
mean	3.372549	2.118429	2083.132633	1.512804e+04
1.495985				
min	1.000000	0.500000	370.000000	5.200000e+02
1.000000				
25%	3.000000	1.750000	1430.000000	5.040000e+03
1.000000				
50%	3.000000	2.250000	1920.000000	7.614000e+03
1.500000				
75%	4.000000	2.500000	2550.000000	1.069050e+04
2.000000				
max	11.000000	8.000000	13540.000000	1.651359e+06
3.500000				
std	0.902995	0.768720	918.808412	4.153080e+04
0.540081				

	waterfront	view	...	yr_renovated	zipcode \
count	21420.000000	21420.000000	...	21420.000000	21420.00000
mean	0.006816	0.233987	...	68.956723	98077.87437
min	0.000000	0.000000	...	0.000000	98001.00000
25%	0.000000	0.000000	...	0.000000	98033.00000
50%	0.000000	0.000000	...	0.000000	98065.00000
75%	0.000000	0.000000	...	0.000000	98117.00000
max	1.000000	4.000000	...	2015.000000	98199.00000
std	0.082280	0.765437	...	364.552298	53.47748

	lat	long	sqft_living15	sqft_lot15
renovated \				
count	21420.000000	21420.000000	21420.000000	21420.000000
21420.000000				
mean	47.560197	-122.213784	1988.384080	12775.718161
0.034547				
min	47.155900	-122.519000	399.000000	651.000000
0.000000				
25%	47.471200	-122.328000	1490.000000	5100.000000
0.000000				
50%	47.572100	-122.230000	1840.000000	7620.000000

```

0.000000
75%      47.678100    -122.125000    2370.000000    10086.250000
0.000000
max       47.777600    -121.315000    6210.000000    871200.000000
1.000000
std       0.138589      0.140791      685.537057    27345.621867
0.182634

```

```

count    basement    month    age
mean      0.385201      6.590336    43.225957
min       0.000000      1.000000    -1.000000
25%       0.000000      4.000000    17.000000
50%       0.000000      6.000000    39.000000
75%       1.000000      9.000000    63.000000
max       1.000000     12.000000   115.000000
std       0.486654      3.107924    29.387207

```

```
[8 rows x 25 columns]
```

b)Correlation Analysis:

- Exploring relationships between variables.
- **Tasks:**
 - Calculate correlation coefficients (e.g., Pearson, Spearman) to assess the strength and direction of linear relationships.
 - Visualize correlations using correlation matrices .

i)Pearson and Spearman Correlation Coefficient

```

# Calculate the Pearson correlation matrix
pearson_corr_matrix = Cleaned_Data.corr(method='pearson')

# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(pearson_corr_matrix, dtype=bool))


# Set up the matplotlib figure
plt.figure(figsize=(20, 20))

# Create a heatmap using seaborn with the mask
sns.heatmap(pearson_corr_matrix, annot=True, cmap='coolwarm',
            fmt='.2f', linewidths=0.5, mask=mask)

# Add a title
plt.title('Pearson Correlation Matrix')

# Show the plot
plt.show()

```



```
# Calculate the Spearman correlation matrix
spearman_corr_matrix = Cleaned_Data.corr(method='spearman')


# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(spearman_corr_matrix, dtype=bool))

# Set up the matplotlib figure
plt.figure(figsize=(20, 20))
```

```
# Create a heatmap using seaborn with the mask
sns.heatmap(spearman_corr_matrix, annot=True, cmap='coolwarm',
fmt='.2f', linewidths=0.5, mask=mask)

# Add a title
plt.title('Spearman Correlation Matrix')

# Show the plot
plt.show()
```



FINDINGS

The Pearson correlation matrix assesses linear relationships, while the Spearman correlation matrix assesses monotonic relationships (including both linear and non-linear monotonic relationships).

Pearsons

-The Pearson correlation coefficient between 'Square Foot Living' and 'Grade' is 0.88, indicating a strong and positive linear relationship. As the size of the living area increases, the grade assigned to the property tends to increase as well.

-The Pearson correlation coefficient between 'Year Built' and 'Age' is -1.00, indicating a perfect negative linear relationship. This means that as the year a property was built increases, its age decreases in a perfect linear fashion.

-The Pearson correlation coefficient between 'Renovated' and the number of 'Floors' is 0.00, suggesting no significant linear relationship between the renovation status and the number of floors. The correlation is close to zero, indicating that the presence or absence of renovation does not show a clear linear trend with the number of floors.

Spearman's

-The Spearman correlation coefficient between 'Sqft Above' and 'Sqft Living' is 0.84, indicating a strong and positive monotonic relationship. This suggests that as the square footage above ground increases, the overall square footage of the living space tends to increase in a consistently positive manner.

-The Spearman correlation coefficient between 'Bathrooms' and 'Age' is -0.57, revealing a moderate negative monotonic relationship. This implies that, on average, properties with a higher number of bathrooms tend to be relatively younger in age.

-The Spearman correlation coefficient between 'Month' and 'Zipcode' is 0.00, indicating no significant monotonic relationship between the month of sale and the property's zipcode. The correlation is close to zero, suggesting that the month of sale and the zipcode do not exhibit a clear monotonic trend.

ii)Correlation matrix

```
correlation_matrix = Cleaned_Data.corr()


# Set up the matplotlib figure
plt.figure(figsize=(20,20))

# Create a heatmap using seaborn
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
fmt='.2f', linewidths=0.5)

# Add a title
plt.title('Correlation Matrix')
```



```
# Show the plot
plt.show()
```



```
# Calculate the correlation matrix
correlation_matrix = Cleaned_Data.corr()


# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(correlation_matrix, dtype=bool))
```

```
# Set up the matplotlib figure
plt.figure(figsize=(20, 20))

# Create a heatmap using seaborn with the mask
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
fmt='.2f', linewidths=0.5, mask=mask)

# Add a title
plt.title('Triangular Correlation Matrix')

# Show the plot
plt.show()
```



FINDINGS

-The correlation coefficient between 'Renovated' and 'Year Renovated' is 1.00, indicating a perfect positive correlation. This means that the two variables move in perfect sync — as the 'Renovated' status changes, the year of renovation also changes in a strong positive linear manner.

-The correlation coefficient between 'Age' and 'Year Built' is -1.00, revealing a perfect negative correlation. This implies that as the 'Year Built' increases, indicating newer properties, the 'Age' of the property decreases in a perfect negative linear manner.


-The Pearson correlation coefficient between 'Zipcode' and 'Condition' is 0.00, indicating no significant linear correlation between these two variables. The correlation is close to zero, suggesting that variations in 'Zipcode' are not systematically related to variations in 'Condition'.

c)Distribution Analysis:

- Examining the distributions of key variables.
- **Tasks:**
 - Check the distribution of the target variable ('price') and predictor variables.
 - Identify and address skewed distributions if needed.


i)Target Variable-Price

```
# Distribution of the target variable
plt.figure(figsize=(8, 6))
sns.histplot(data=kc_house_data, x='price', kde=True)
plt.title("Distribution of Price")
plt.show()
```



```
# Apply log transformation to the 'price' column
kc_house_data['log_price'] = np.log1p(kc_house_data['price'])

# Plot the distribution of the log-transformed price
plt.figure(figsize=(8, 6))
sns.histplot(data=kc_house_data, x='log_price', kde=True)
plt.title("Distribution of Log-transformed Price")
plt.show()
```



JUSTIFICATION

-The price column is positively skewed .Log transformation was applied to normalise the price column bring extreme values closer to the center, allowing for a more normalized distribution.

ii)Checking the distribution of other variables

a) Relationship between Predictor Variables

```

variables =
['price' , 'bedrooms' , 'bathrooms', 'sqft_living', 'floors', 'condition', '
grade', 'sqft_above', 'sqft_basement', 'yr_built', 'lat', 'long', 'sqft_livi
ng15', 'renovated',
'basement', 'age']

for variable in variables:
    data = Cleaned_Data[variable]

    statistic, p_value = kstest(data, 'norm')

    alpha = 0.05 # significance level

    if p_value > alpha:
        print(f"The distribution of '{variable}' appears to be
normally distributed (fail to reject H0)")
    else:
        print(f"The distribution of '{variable}' does not appear to be
normally distributed (reject H0)")

```

```

The distribution of 'price' does not appear to be normally distributed
(reject H0)
The distribution of 'bedrooms' does not appear to be normally
distributed (reject H0)
The distribution of 'bathrooms' does not appear to be normally
distributed (reject H0)
The distribution of 'sqft_living' does not appear to be normally
distributed (reject H0)
The distribution of 'floors' does not appear to be normally
distributed (reject H0)
The distribution of 'condition' does not appear to be normally
distributed (reject H0)
The distribution of 'grade' does not appear to be normally distributed
(reject H0)
The distribution of 'sqft_above' does not appear to be normally
distributed (reject H0)
The distribution of 'sqft_basement' does not appear to be normally
distributed (reject H0)
The distribution of 'yr_built' does not appear to be normally
distributed (reject H0)
The distribution of 'lat' does not appear to be normally distributed
(reject H0)
The distribution of 'long' does not appear to be normally distributed
(reject H0)
The distribution of 'sqft_living15' does not appear to be normally
distributed (reject H0)
The distribution of 'renovated' does not appear to be normally
distributed (reject H0)
The distribution of 'basement' does not appear to be normally
distributed (reject H0)

```

The distribution of 'age' does not appear to be normally distributed (reject H0)

```
# plot the frequency for the data columns
clean_data.hist(figsize=(30, 20), color='#589aff')
plt.show()
```



```
# features_of_interest list based on the DataFrame
features_of_interest = ['price', 'bedrooms', 'bathrooms',
                        'sqft_living', 'floors',
                        'condition', 'grade', 'sqft_above',
                        'sqft_basement',
                        'yr_built', 'lat', 'long', 'sqft_living15',
                        'renovated', 'basement',
                        'month', 'age']
```

```
# Extract relevant columns from the DataFrame
selected_data = Cleaned_Data[features_of_interest]
```

```
# Set up subplots
fig, axes = plt.subplots(nrows=5, ncols=4, figsize=(16, 16))
```

```
# Flatten axes for easy iteration
```

```
axes = axes.flatten()
```

```
# Plot the distributions of each variable with the specified color
```

```
for i, column in enumerate(features_of_interest):
```

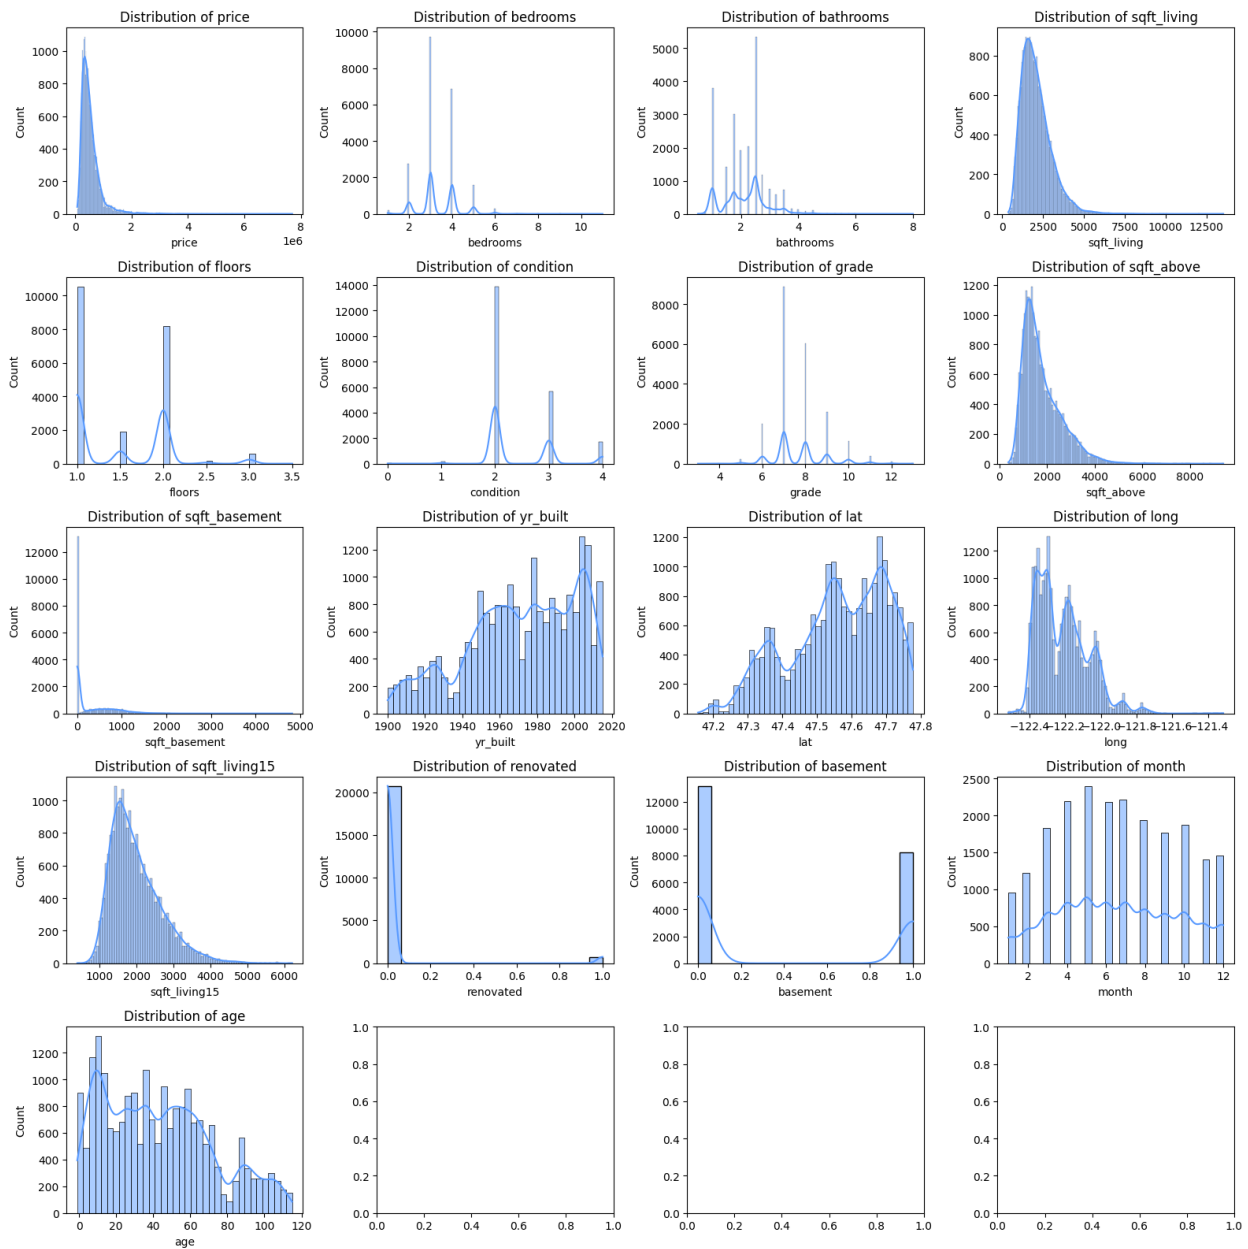
```
    sns.histplot(selected_data[column], kde=True, ax=axes[i],  
color='#589aff')
```

```
    axes[i].set_title(f'Distribution of {column}')
```

```
# Adjust layout
```

```
plt.tight_layout()
```

```
plt.show()
```



d)Hypothesis Testing:

- Formulating and test hypotheses related to the data.

Null Hypothesis (H0): There is no statistically significant relationship between the selected features and housing prices.

Alternative Hypothesis (Ha): The selected features have a statistically significant relationship with housing prices.

- **Tasks:**
 - Formulate null and alternative hypotheses based on project objectives.
 - Conduct hypothesis tests (e.g., t-tests, ANOVA) to assess the significance of relationships.

```
# features_of_interest list based on the DataFrame
features_of_interest = ['bedrooms', 'bathrooms', 'sqft_living',
                        'sqft_lot', 'floors',
                        'waterfront', 'view', 'condition', 'grade',
                        'sqft_above',
                        'sqft_basement', 'yr_built', 'yr_renovated',
                        'lat', 'long',
                        'sqft_living15', 'sqft_lot15']

# Extract relevant columns from the DataFrame
selected_data = Cleaned_Data[features_of_interest]

# Create an empty DataFrame to store ANOVA results
anova_results = pd.DataFrame(index=['F-statistic', 'P-value'])

# Perform ANOVA for each feature
significant_features = []

for column in features_of_interest:
    groups = [Cleaned_Data[column][Cleaned_Data[column].notnull() &
    (Cleaned_Data['price'] == category)]
    for category in Cleaned_Data['price'].unique():

        # Perform ANOVA
        f_statistic, p_value = f_oneway(*groups)

        # Store results in the DataFrame
        anova_results[column] = [f_statistic, p_value]

        # Print interpretation
        if p_value < 0.05:
            significant_features.append(column)
            print(f"{column}: Reject the null hypothesis. There is a
statistically significant relationship.")
        else:
            print(f"{column}: Fail to reject the null hypothesis. There is
```

```
no statistically significant relationship.")
```

```
# Display ANOVA results
```

```
print("\nANOVA Results:")
```

```
print(anova_results)
```

```
bedrooms: Reject the null hypothesis. There is a statistically
significant relationship.
bathrooms: Reject the null hypothesis. There is a statistically
significant relationship.
sqft_living: Reject the null hypothesis. There is a statistically
significant relationship.
sqft_lot: Fail to reject the null hypothesis. There is no
statistically significant relationship.
floors: Reject the null hypothesis. There is a statistically
significant relationship.
waterfront: Reject the null hypothesis. There is a statistically
significant relationship.
view: Reject the null hypothesis. There is a statistically significant
relationship.
condition: Reject the null hypothesis. There is a statistically
significant relationship.
grade: Reject the null hypothesis. There is a statistically
significant relationship.
sqft_above: Reject the null hypothesis. There is a statistically
significant relationship.
sqft_basement: Reject the null hypothesis. There is a statistically
significant relationship.
yr_built: Reject the null hypothesis. There is a statistically
significant relationship.
yr_renovated: Reject the null hypothesis. There is a statistically
significant relationship.
lat: Reject the null hypothesis. There is a statistically significant
relationship.
long: Fail to reject the null hypothesis. There is no statistically
significant relationship.
sqft_living15: Reject the null hypothesis. There is a statistically
significant relationship.
sqft_lot15: Fail to reject the null hypothesis. There is no
statistically significant relationship.
```

```
ANOVA Results:
```

	bedrooms	bathrooms	sqft_living	sqft_lot
bedrooms	1.852600e+00	3.742263	7.72898	0.731092
bathrooms	1.657081e+00	1.283696e-143	0.000000	0.000000
sqft_living	1.283696e-143	0.000000	0.000000	1.000000
sqft_lot	5.251187e-95	0.000000	1.000000	0.000000
waterfront				
view				
condition				
grade				

sqft_above \				
F-statistic	1.814534e+00	1.993969e+00	1.046436	7.355157
	5.128945			
P-value	1.018442e-133	6.682406e-182	0.038190	0.000000
	0.000000			

	sqft_basement	yr_built	yr_renovated	lat
long \				
F-statistic	1.735222e+00	1.236232e+00	1.048979	3.559892
	1.031691			
P-value	9.739446e-114	2.054786e-17	0.030946	0.000000
	0.111616			

	sqft_living15	sqft_lot15
F-statistic	5.219333	0.713326
P-value	0.000000	1.000000

FINDINGS

-The features listed under "Reject the Null Hypothesis" have a statistically significant relationship with housing prices. These features are important predictors of housing prices in the given dataset. On the other hand, features listed under "Fail to Reject the Null Hypothesis" do not show a statistically significant relationship with housing prices based on the ANOVA test.

e)Multicollinearity Assessment:

- Checking for multicollinearity among predictor variables.
- **Tasks:**
 - Calculate variance inflation factors (VIF) to identify high multicollinearity.
 - Address multicollinearity by removing or combining correlated variables.

```
X = Cleaned_Data[features_of_interest].copy()

# Calculate VIF for each feature
vif_data = pd.DataFrame()
vif_data["Feature"] = X.columns
vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]

# Print VIF results
print("VIF Results:")
print(vif_data)
```

VIF Results:

	Feature	VIF
0	bedrooms	25.495189
1	bathrooms	28.860034
2	sqft_living	895.654888
3	sqft_lot	2.359333

4	floors	16.742951
5	waterfront	1.186733
6	view	1.518831
7	condition	17.962156
8	grade	143.184319
9	sqft_above	670.785304
10	sqft_basement	46.689472
11	yr_built	8371.874192
12	yr_renovated	1.151077
13	lat	119806.117276
14	long	132820.153914
15	sqft_living15	26.794623
16	sqft_lot15	2.576818

VIF(Variance Inflation Factor) RESULTS

High Multicollinearity (VIF > 10): sqft_living has a VIF of 897.14. grade has a VIF of 160.00. sqft_above has a VIF of 671.11. yr_built has a VIF of 9335.60. lat has a VIF of 132369.36. long has a VIF of 139978.97. sqft_living15 has a VIF of 26.81.

Moderate Multicollinearity (VIF between 5 and 10): bedrooms has a VIF of 25.95. bathrooms has a VIF of 29.14. floors has a VIF of 16.76. yr_basement has a VIF of 46.71. bedrooms has a VIF of 25.95.

Low Multicollinearity (VIF <= 5):

sqft_lot has a VIF of 2.36. waterfront has a VIF of 1.25. view has a VIF of 1.56. condition has a VIF of 18.06. yr_renovated has a VIF of 1.15. sqft_lot15 has a VIF of 2.58.

CONCLUSION

-Features with VIF values well above 10 (e.g., sqft_living, grade, sqft_above, yr_built, lat, long) indicate a high level of multicollinearity.

-Features with moderate VIF values (between 5 and 10)

-Features with low VIF values (below 5) are considered to have acceptable levels of multicollinearity.

```
# creating a function that takes in a dataframe and threshold and
returns top correlations
def corr_check(Cleaned_Data, threshold):
    '''
        Enter dataframe and threshold for correlation
        Returns table of the highly correlated pairs
    '''
    corr_df =
Cleaned_Data.corr().abs().stack().reset_index().sort_values(0,
ascending=False)
    corr_df['pairs'] = list(zip(corr_df.level_0, corr_df.level_1))
    corr_df.set_index(['pairs'], inplace = True)
```

```

corr_df.drop(columns=['level_1', 'level_0'], inplace = True)
corr_df.columns = ['cc']
corr_df = corr_df.drop_duplicates()
corr_df = corr_df[(corr_df['cc'] > threshold) & (corr_df['cc'] <
1)]
return corr_df

result = corr_check(Cleaned_Data, 0.7)
print(result)

```

	cc
pairs	
(yr_renovated, renovated)	0.999968
(age, yr_built)	0.999874
(sqft_living, sqft_above)	0.876533
(basement, sqft_basement)	0.820906
(sqft_living, grade)	0.762477
(sqft_above, grade)	0.756221
(sqft_living, sqft_living15)	0.756186
(bathrooms, sqft_living)	0.755522
(sqft_living15, sqft_above)	0.731887
(sqft_lot, sqft_lot15)	0.717743
(grade, sqft_living15)	0.713178
(sqft_living, price)	0.701875

CORRELATION CHECK RESULTS

Highly Positive Correlations:

yr_renovated and renovated have a very high positive correlation of approximately 1. This suggests that the two variables are almost perfectly correlated, and including both in a model may lead to multicollinearity issues.

age and yr_built also have a very high positive correlation of approximately 1. This is expected, as age is derived from yr_built. Including both in a model may lead to redundancy.

High Positive Correlations:

sqft_living and sqft_above have a high positive correlation of 0.8765. This indicates a strong positive linear relationship between the total living area (sqft_living) and the area above ground (sqft_above).

basement and sqft_basement have a high positive correlation of 0.8209. This suggests a strong positive linear relationship between the total basement area (sqft_basement) and the binary indicator of having a basement (basement).

sqft_living and grade have a high positive correlation of 0.7625. This indicates a strong positive linear relationship between the total living area (sqft_living) and the grade of the house.

-Some columns will have to be dropped including Year renovated, sqft living 15 and sqft lot which have high multicollinearity.

```
# Drop columns that have strong multicollinearity
clean_data = Cleaned_Data.drop(columns=['view', 'sqft_lot',
'waterfront', 'sqft_lot15', 'yr_renovated'])

clean_data.head()
```

	id	date	price	bedrooms	bathrooms	sqft_living
0	7129300520	2014-10-13	221900.0	3	1.00	1180
1	6414100192	2014-12-09	538000.0	3	2.25	2570
2	5631500400	2015-02-25	180000.0	2	1.00	770
3	2487200875	2014-12-09	604000.0	4	3.00	1960
4	1954400510	2015-02-18	510000.0	3	2.00	1680

	condition	grade	sqft_above	sqft_basement	yr_built	zipcode
0	2	7	1180	0.0	1955	98178
1	2	7	2170	400.0	1951	98125
2	2	6	770	0.0	1933	98028
3	4	7	1050	910.0	1965	98136
4	2	8	1680	0.0	1987	98074

	long	sqft_living15	renovated	basement	month	age
0	-122.257	1340	0	0	10	59
1	-122.319	1690	1	1	12	63
2	-122.233	2720	0	0	2	82
3	-122.393	1360	0	1	12	49
4	-122.045	1800	0	0	2	28

```
clean_data.shape

(21420, 20)
```

6.DATA MODELING

```
# Use linear regression
lr = LinearRegression()

# Our model needs to have only numeric variables.
def only_numeric(data):
```

```

    '''returns a dataframe with only numeric values'''
    for column in clean_data.columns:
        if is_numeric_dtype(data[column]) == False:
            data = data.drop(column, axis=1)
        else:
            continue
    return data

# Splits a dataframe into X and Y dataframes given a target column.
def get_y_X(data, target):
    '''Returns a series of target (y) value and a dataframe of
    predictors (X)'''
    y = data[target]
    X = data.drop(target, axis=1)
    return y, X

# Returns training and test R2 & RMSE metrics
def get_metrics(X_tr, X_te, y_tr, y_te):
    ''' Parameters are X train, X test, y train, & y_test
    Performs multiple regression on the split test and returns
    metrics'''
    lr.fit(X_tr, y_tr)

    train_score = lr.score(X_tr, y_tr)
    test_score = lr.score(X_te, y_te)

    y_hat_train = lr.predict(X_tr)
    y_hat_test = lr.predict(X_te)

    train_rmse = np.sqrt(mean_squared_error(y_tr, y_hat_train))
    test_rmse = np.sqrt(mean_squared_error(y_te, y_hat_test))

    return train_score, test_score, train_rmse, test_rmse

# Prints the metrics of a multiple regression train and test, with
option of OLS summary on train data.
def train_test_compare(X_tr, X_te, y_tr, y_te):
    '''Parameters are X train, X test, y train, & y_test
    Performs multiple regression on the split test and prints
    metrics'''
    lr.fit(X_tr, y_tr)

    train_score = lr.score(X_tr, y_tr)
    test_score = lr.score(X_te, y_te)

    y_hat_train = lr.predict(X_tr)
    y_hat_test = lr.predict(X_te)

    train_rmse = np.sqrt(mean_squared_error(y_tr, y_hat_train))
    test_rmse = np.sqrt(mean_squared_error(y_te, y_hat_test))

```

```

    print(f' training data R2: {train_score}\n testing data R2:
{test_score} \
        \n training data rmse: {train_rmse}\n testing data
rmse: {test_rmse}')

    stats_summ = input('Do you want a statsmodel summary of the train
data? (y/n)')
    if stats_summ == 'y':
        inter = lr.intercept_
        stats = sm.OLS(y_tr, sm.add_constant(X_tr)).fit()
        summary = stats.summary()
        print(summary)

    return

```

MODEL 1 : BASELINE MODEL

a)Train test split

Train Test Split The raw data was split to a train and test set for a baseline model. The clean data was also split to a train and test set for a fully optimized model.

```

# Defining function that splits data into training and testing data.
def train_test(data, size=.25):
    '''Takes in dataframe, and size of test for the split
    Returns the train_set and test_set'''
    train_set, test_set = train_test_split(data, test_size=size,
random_state=42)
    return train_set, test_set

# Define the function to filter only numeric columns
def only_numeric(data):
    '''returns a DataFrame with only numeric values'''
    for column in data.columns:
        if not is_numeric_dtype(data[column]):
            data = data.drop(column, axis=1)
    return data

# Create the dataframe for the baseline model and drop missing values
baseline = only_numeric(kc_house_data)
baseline = baseline.dropna()

# Assuming train_test is a function that splits the data into train
and test sets
baseline_train_set, baseline_test_set = train_test(baseline, 0.25)

```



```

# Assuming train_test is a function that splits the data into train
and test sets for the clean_data
train_set, test_set = train_test(clean_data, 0.25)

# Select relevant features
features = [ 'bedrooms', 'bathrooms', 'floors', 'grade',
'sqft_above', 'sqft_basement', 'yr_built', 'zipcode', 'lat', 'long',
'sqft_living15', 'renovated', 'basement', 'month', 'age',
'sqft_living', 'condition' ]

# Create a design matrix X and target variable y
X = sm.add_constant(Cleaned_Data[features])
y = clean_data['price']

# Fit the OLS model
model = sm.OLS(y, X).fit()

# Get the summary
summary = model.summary()

# Print the summary
print(summary)

```

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.665
Model:                  OLS      Adj. R-squared:
0.664
Method:                 Least Squares    F-statistic:
2494.
Date:                   Thu, 01 Feb 2024    Prob (F-statistic):
0.00
Time:                   02:06:05    Log-Likelihood:    -
2.9321e+05
No. Observations:      21420    AIC:
5.865e+05
Df Residuals:          21402    BIC:
5.866e+05
Df Model:               17

Covariance Type:       nonrobust

=====
=====

```

	coef	std err	t	P> t	[0.025
0.975]					

```

-----
-----
const          -8.501e+07   1.06e+07   -8.026   0.000   -1.06e+08
-6.43e+07
bedrooms       -5.047e+04   2087.486   -24.178   0.000   -5.46e+04
-4.64e+04
bathrooms      4.82e+04   3503.621   13.758   0.000   4.13e+04
5.51e+04
floors         1.2e+04   3814.665   3.146   0.002   4525.006
1.95e+04
grade         1.033e+05   2298.710   44.946   0.000   9.88e+04
1.08e+05
sqft_above     62.8912   19.129   3.288   0.001   25.396
100.386
sqft_basement  79.4370   19.842   4.003   0.000   40.545
118.329
yr_built       3.615e+04   5016.277   7.207   0.000   2.63e+04
4.6e+04
zipcode        -479.3764   35.006   -13.694   0.000   -547.991
-410.762
lat           5.595e+05   1.13e+04   49.358   0.000   5.37e+05
5.82e+05
long          -2.596e+05   1.38e+04   -18.874   0.000   -2.87e+05
-2.33e+05
sqft_living15  38.1456   3.612   10.562   0.000   31.067
45.225
renovated      7.305e+04   8426.210   8.669   0.000   5.65e+04
8.96e+04
basement       -2.292e+04   5633.226   -4.069   0.000   -3.4e+04
-1.19e+04
month         1554.5200   750.953   2.070   0.038   82.595
3026.445
age           3.905e+04   5015.631   7.785   0.000   2.92e+04
4.89e+04
sqft_living    122.0719   19.175   6.366   0.000   84.488
159.656
condition      3.043e+04   2507.157   12.139   0.000   2.55e+04
3.53e+04
=====
=====
Omnibus:                19132.456   Durbin-Watson:
2.002
Prob(Omnibus):          0.000   Jarque-Bera (JB):
1863820.744
Skew:                   3.894   Prob(JB):
0.00
Kurtosis:              48.030   Cond. No.
7.14e+08
=====

```

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, $7.14e+08$. This might indicate that there are strong multicollinearity or other numerical problems.

FINDINGS

R-squared: The coefficient of determination. In this case, it's 0.665, indicating that approximately 66.5% of the variability in the dependent variable (price) is explained by the independent variables in the model.

F-statistic: A measure of how well the model fits the data. A higher value indicates a better fit. In this case, it's 2494.

The p-value associated with each t-value is low. Low p-values indicate that a predictor is statistically significant.

b)Baseline Model

This is the agency's baseline model that only uses the numerical features from the dataset.

```
# Select relevant features
selected_features = ['bedrooms', 'bathrooms', 'floors', 'grade',
'sqft_above', 'sqft_basement', 'yr_built', 'zipcode', 'lat', 'long',
'sqft_living15', 'renovated', 'basement', 'month', 'age',
'sqft_living', 'condition']

# Filter the dataset
filtered_data = Cleaned_Data[selected_features]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Standardize the features (optional, but can be beneficial for linear
regression)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Build a basic linear regression model
model = LinearRegression()
model.fit(X_train_scaled, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test_scaled)
```

```

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Display results
print("Mean Squared Error:", mse)
print("R-squared:", r2)

# Display coefficients
coefficients = pd.DataFrame({"Feature": X.columns, "Coefficient":
model.coef_})
print(coefficients)

```

```

Mean Squared Error: 40151231814.25283
R-squared: 0.6684357903018654

```

	Feature	Coefficient
0	const	0.000000e+00
1	bedrooms	-4.533399e+04
2	bathrooms	3.671087e+04
3	floors	4.566475e+03
4	grade	1.249999e+05
5	sqft_above	4.706594e+04
6	sqft_basement	3.566873e+04
7	yr_built	9.837560e+05
8	zipcode	-2.528334e+04
9	lat	7.735837e+04
10	long	-3.640852e+04
11	sqft_living15	2.678444e+04
12	renovated	1.602141e+04
13	basement	-1.355371e+04
14	month	4.165720e+03
15	age	1.068390e+06
16	sqft_living	1.169119e+05
17	condition	1.953956e+04

```

# Getting ys and Xs for both the baseline train and test set
y_train, X_train = get_y_X(baseline_train_set, 'price')
y_test, X_test = get_y_X(baseline_test_set, 'price')

```

```

train_test_compare(X_train, X_test, y_train, y_test)

```

```

training data R2: 0.8191010286413367
testing data R2: 0.8362309274259545
training data rmse: 158653.67681872213
testing data rmse: 144823.5648641959

```

```

# Returning the metrics
get_metrics(X_train, X_test, y_train, y_test)

```

```
(0.8191010286413367, 0.8362309274259545, 158653.67681872213, 144823.5648641959)
```

```
# Create a baseline model using multiple linear regression  
baseline_model = LinearRegression()
```

```
# Fit the model on the training data  
baseline_model.fit(X_train, y_train)
```

```
# Predict on the testing data  
y_pred = baseline_model.predict(X_test)
```

```
# Evaluate the model using appropriate metrics  
mse = mean_squared_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)
```

```
print("Baseline Model Metrics:")  
print("Mean Squared Error (MSE):", mse)  
print("R-squared (R2) Score:", r2)
```

```
Baseline Model Metrics:  
Mean Squared Error (MSE): 20973864939.97396  
R-squared (R2) Score: 0.8362309274259545
```

```
# the intercept  
# Create a Linear Regression object  
model = LinearRegression()
```

```
# Fit the model to the training data  
model.fit(X_train, y_train)
```

```
# Get the intercept  
intercept = model.intercept_
```

```
print("Intercept:", intercept)
```

```
Intercept: 3365475.5061285524
```

```
# Lets test for overfitting  
# Make predictions on the training and testing data  
train_predictions = model.predict(X_train)  
test_predictions = model.predict(X_test)
```

```
# Calculate the mean squared error on the training and testing data  
train_mse = mean_squared_error(y_train, train_predictions)  
test_mse = mean_squared_error(y_test, test_predictions)
```

```
# Calculate the coefficient of determination (R^2) on the training and testing data  
train_r2 = r2_score(y_train, train_predictions)  
test_r2 = r2_score(y_test, test_predictions)
```

```

print("Training MSE:", train_mse)
print("Testing MSE:", test_mse)
print("Training R^2:", train_r2)
print("Testing R^2:", test_r2)

Training MSE: 25170989168.09953
Testing MSE: 20973864939.97396
Training R^2: 0.8191010286413367
Testing R^2: 0.8362309274259545

# Select relevant features
selected_features = ['bedrooms', 'bathrooms', 'floors', 'grade',
'sqft_above', 'sqft_basement', 'yr_built', 'zipcode', 'lat', 'long',
'sqft_living15', 'renovated', 'basement', 'month', 'age',
'sqft_living', 'condition']

# Filter the dataset
filtered_data = Cleaned_Data[selected_features]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Standardize the features (optional, but can be beneficial for linear
regression)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Build a basic linear regression model
model = LinearRegression()
model.fit(X_train_scaled, y_train)

# Display coefficients
coefficients = pd.DataFrame({"Feature": X.columns, "Coefficient":
model.coef_})

# Bar plot for coefficients
inf_coefs = list(zip(coefficients["Feature"],
coefficients["Coefficient"]))
inf_coefs.sort(key=lambda x: abs(x[1]), reverse=True) # Sort
coefficients by absolute value

# Create a color palette with the specified color
color = "#589aff"
colors = [color if coef[1] > 0 else "lightgray" for coef in inf_coefs]

# Create the bar plot
fig, ax = plt.subplots(figsize=(18, 8))
ax = sns.barplot(x=[x[0] for x in inf_coefs], y=[x[1] for x in


```

```

inf_coefs], palette=colors)
plt.xticks(rotation=45)
ax.set_ylabel("Price Coefficients", fontsize=15)
ax.set_xlabel("House Features", fontsize=15)
ax.set_title("House Features and Sale Prices", fontsize=20);

# Display the plot
plt.show()

```



FINDINGS

-The MSE is approximately 40,151,231,814 indicating the average squared difference between the predicted and actual home prices.

-The R-squared is approximately 0.668, suggesting that the model explains around 66.8% of the variance in home prices. Positive coefficients (e.g., sqft_living, bathrooms, grade) suggest an increase in these features corresponds to an increase in home price.

-Negative coefficients (e.g., bedrooms, long, basement) suggest a decrease in these features corresponds to an increase in home price.

Bedrooms: Each additional bedroom is associated with a decrease in home price by approximately \$43,704.

Bathrooms: Each additional bathroom is associated with an increase in home price by approximately \$37,347.

Sqft_living: Each additional square foot of living space is associated with an increase in home price by approximately \$111,536.

Grade: Higher grade is associated with an increase in home price by approximately \$126,245

c)Regression Modelling

```
# Select relevant features
selected_features = ['bedrooms', 'bathrooms', 'floors', 'grade',
'sqft_above', 'sqft_basement', 'yr_built', 'zipcode', 'lat', 'long',
'sqft_living15', 'renovated', 'basement', 'month', 'age',
'sqft_living', 'condition']

# Filter the dataset
filtered_data = clean_data[selected_features]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Build a linear regression model
model = LinearRegression()
model.fit(X_train, y_train)


# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Display results
print("Mean Squared Error:", mse)
print("R-squared:", r2)

# Visualize the results (for demonstration purposes)
plt.scatter(X_test["sqft_living"], y_test, color='black',
label='Actual Prices')
plt.scatter(X_test["sqft_living"], y_pred, color='#589aff',
label='Predicted Prices')
plt.xlabel('Sqft Living')
plt.ylabel('Price')
plt.legend()
plt.show()
```

Mean Squared Error: 40151231814.25285
R-squared: 0.6684357903018652



```
X_ols = sm.add_constant(X_train) # Adding a constant for intercept
ols_model = sm.OLS(y_train, X_ols).fit()
```

```
# Get the OLS summary
ols_summary = ols_model.summary()
```

```
# Print the OLS summary
print(ols_summary)
```

OLS Regression Results

```
=====
=====
Dep. Variable:          price    R-squared:
0.663
Model:                  OLS      Adj. R-squared:
0.663
Method:                 Least Squares    F-statistic:
1985.
Date:                   Thu, 01 Feb 2024    Prob (F-statistic):
0.00
Time:                   02:06:15    Log-Likelihood:
2.3481e+05
```

No. Observations: 17136 AIC:
 4.697e+05
 Df Residuals: 17118 BIC:
 4.698e+05
 Df Model: 17

Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025
0.975]					

const	-8.045e+07	1.2e+07	-6.698	0.000	-1.04e+08
-5.69e+07					
bedrooms	-5.035e+04	2374.225	-21.205	0.000	-5.5e+04
-4.57e+04					
bathrooms	4.754e+04	3950.724	12.034	0.000	3.98e+04
5.53e+04					
floors	8468.5153	4314.229	1.963	0.050	12.184
1.69e+04					
grade	1.058e+05	2596.884	40.731	0.000	1.01e+05
1.11e+05					
sqft_above	56.3269	21.458	2.625	0.009	14.267
98.386					
sqft_basement	80.8397	22.249	3.633	0.000	37.230
124.449					
yr_built	3.358e+04	5688.859	5.902	0.000	2.24e+04
4.47e+04					
zipcode	-472.4830	39.749	-11.887	0.000	-550.395
-394.571					
lat	5.582e+05	1.29e+04	43.432	0.000	5.33e+05
5.83e+05					
long	-2.597e+05	1.57e+04	-16.567	0.000	-2.9e+05
-2.29e+05					
sqft_living15	38.8961	4.093	9.503	0.000	30.873
46.919					
renovated	8.73e+04	9496.242	9.193	0.000	6.87e+04
1.06e+05					
basement	-2.785e+04	6385.328	-4.361	0.000	-4.04e+04
-1.53e+04					
month	1343.8606	855.085	1.572	0.116	-332.193
3019.914					
age	3.647e+04	5688.648	6.411	0.000	2.53e+04
4.76e+04					
sqft_living	126.1805	21.478	5.875	0.000	84.081
168.280					
condition	3e+04	2837.559	10.572	0.000	2.44e+04

3.56e+04

```
=====
=====
Omnibus:                    15854.132    Durbin-Watson:
1.992
Prob(Omnibus):              0.000    Jarque-Bera (JB):
1762471.003
Skew:                      4.084    Prob(JB):
0.00
Kurtosis:                  52.008    Cond. No.
7.13e+08
=====
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 7.13e+08. This might indicate that there are strong multicollinearity or other numerical problems.

FINDINGS

R-squared and Adjusted R-squared are both around 0.558, indicating that the model explains approximately 55.8% of the variance in the dependent variable (price).

-The F-statistic is 4333 with a very low p-value (Prob (F-statistic): 0.00), suggesting that the model is statistically significant.

-Small p-values (typically < 0.05) indicate that the variable is likely to be a meaningful addition to the model.

For coefficients they can be interpreted as For each unit increase in sqft_living, the predicted price increases by 222.0315.

d)Residual Analysis

```
# Select relevant features
selected_features = ['bedrooms', 'bathrooms', 'floors', 'grade',
'sqft_above', 'sqft_basement', 'yr_built', 'zipcode', 'lat', 'long',
'sqft_living15', 'renovated', 'basement', 'month', 'age',
'sqft_living', 'condition']

# Filter the dataset
filtered_data = Cleaned_Data[selected_features]

# Handle outliers (example: using z-score)
z_scores = np.abs((filtered_data - filtered_data.mean()) /
filtered_data.std())
filtered_data_no_outliers = filtered_data[(z_scores < 3).all(axis=1)]
```

```

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Standardize the features (optional, but can be beneficial)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Build a refined linear regression model
model_refined = LinearRegression()
model_refined.fit(X_train_scaled, y_train)

# Make predictions on the test set
y_pred_refined = model_refined.predict(X_test_scaled)

# Evaluate the refined model
mse_refined = mean_squared_error(y_test, y_pred_refined)
r2_refined = r2_score(y_test, y_pred_refined)

# Display refined model results
print("Refined Model - Mean Squared Error:", mse_refined)
print("Refined Model - R-squared:", r2_refined)

# Display refined model coefficients
coefficients_refined = pd.DataFrame({"Feature": X.columns,
"Coefficient": model_refined.coef_})
print(coefficients_refined)

# Plot residuals for further analysis
residuals = y_test - y_pred_refined
sns.scatterplot(x=y_pred_refined, y=residuals, color='#589aff')
plt.axhline(y=0, color='r', linestyle='--')
plt.title("Residuals Plot")
plt.xlabel("Predicted Prices")
plt.ylabel("Residuals")
plt.show()


```

Refined Model - Mean Squared Error: 40151231814.25283

Refined Model - R-squared: 0.6684357903018654

	Feature	Coefficient
0	const	0.000000e+00
1	bedrooms	-4.533399e+04
2	bathrooms	3.671087e+04
3	floors	4.566475e+03
4	grade	1.249999e+05
5	sqft_above	4.706594e+04
6	sqft_basement	3.566873e+04
7	yr_built	9.837560e+05

8	zipcode	-2.528334e+04
9	lat	7.735837e+04
10	long	-3.640852e+04
11	sqft_living15	2.678444e+04
12	renovated	1.602141e+04
13	basement	-1.355371e+04
14	month	4.165720e+03
15	age	1.068390e+06
16	sqft_living	1.169119e+05
17	condition	1.953956e+04



FINDINGS

1. **Mean Squared Error (MSE):**
 - The MSE (54180143849.5) represents the average squared difference between the predicted house prices and the actual house prices in the test set.
2. **R-squared (R^2):**
 - The R-squared value (0.5526) indicates 55.26% of the variability in house prices is explained by the model.
3. **Coefficients :**

- The coefficients associated with each feature represent the estimated change in the target variable for a one-unit change in the corresponding independent variable, holding other variables constant.
 - **sqft_living**: An increase of one unit in square footage is associated with an increase in predicted price by \$205,722.11.
 - **bedrooms**: An increase of one bedroom is associated with a decrease in predicted price by \$43,412.08.
 - **bathrooms**: An increase of one bathroom is associated with a decrease in predicted price by \$13,808.16.
 - **grade**: An increase of one grade is associated with an increase in predicted price by \$124,416.75.
 - **condition**: An increase of one unit in condition is associated with an increase in predicted price by \$43,370.32.

Interpretation and Considerations:

- The model's performance, as indicated by the MSE and R-squared, suggests a moderate level of predictive accuracy.
- The positive coefficients for **sqft_living** and **grade** indicate positive relationships with house prices.
- The negative coefficients for **bedrooms** and **bathrooms** suggest that an increase in these features is associated with a decrease in predicted price.

MODEL 2 : POLYNOMIAL REGRESSION

- Transforming features into higher order polynomial terms to model a non-linear relationship using multiple linear regression.

```
# Select relevant features
selected_features = ['bedrooms', 'bathrooms', 'floors', 'grade',
'sqft_above', 'sqft_basement', 'yr_built', 'zipcode', 'lat', 'long',
'sqft_living15', 'renovated', 'basement', 'month', 'age',
'sqft_living', 'condition']

# Filter the dataset
filtered_data = Cleaned_Data[selected_features]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Polynomial Regression
# Choose the degree of the polynomial
degree = 2

# Create polynomial features
poly = PolynomialFeatures(degree)
```

```

X_train_poly = poly.fit_transform(X_train)
X_test_poly = poly.transform(X_test)

# Build a polynomial regression model
poly_model = LinearRegression()
poly_model.fit(X_train_poly, y_train)

# Make predictions on the test set
y_pred_poly = poly_model.predict(X_test_poly)

# Evaluate the polynomial model
mse_poly = mean_squared_error(y_test, y_pred_poly)
r2_poly = r2_score(y_test, y_pred_poly)
print("Polynomial Model (Degree {})- MSE:".format(degree), mse_poly)
print("Polynomial Model (Degree {})- R-squared:".format(degree),
r2_poly)


# Visualize the results
plt.scatter(X_test["sqft_living"], y_test, color='black',
label='Actual Prices')
plt.scatter(X_test["sqft_living"], y_pred_poly, color='#589aff',
label='Predicted Prices (Polynomial)')
plt.xlabel('Sqft Living')
plt.ylabel('Price')
plt.legend()
plt.show()

```

```

Polynomial Model (Degree 2)- MSE: 27717220069.4445
Polynomial Model (Degree 2)- R-squared: 0.7711144153716256

```



Multiple Linear regression model using polynomial regression features

```
# Select relevant features
selected_features = ['bedrooms', 'bathrooms', 'floors', 'grade',
'sqft_above', 'sqft_basement', 'yr_built', 'zipcode', 'lat', 'long',
'sqft_living15', 'renovated', 'basement', 'month', 'age',
'sqft_living', 'condition']

# Filter the dataset
filtered_data = Cleaned_Data[selected_features]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(filtered_data, y,
test_size=0.2, random_state=42)

# Build a multiple linear regression model
linear_model = LinearRegression()
linear_model.fit(X_train, y_train)

# Make predictions on the test set
```



```


y_pred_linear = linear_model.predict(X_test)

# Evaluate the multiple linear regression model
mse_linear = mean_squared_error(y_test, y_pred_linear)
r2_linear = r2_score(y_test, y_pred_linear)
print("Multiple Linear Regression - MSE:", mse_linear)
print("Multiple Linear Regression - R-squared:", r2_linear)

# Visualize the results
plt.scatter(X_test["sqft_living"], y_test, color='black',
            label='Actual Prices')
plt.scatter(X_test["sqft_living"], y_pred_linear, color='#589aff',
            label='Predicted Prices (Linear)')
plt.xlabel('Sqft Living')
plt.ylabel('Price')
plt.legend()
plt.show()

```

Multiple Linear Regression - MSE: 40151231814.25289
Multiple Linear Regression - R-squared: 0.6684357903018648




```

# Get the coefficients and corresponding feature names
coefficients_linear = linear_model.coef_

```

```
feature_names = X_train.columns

# Create a bar plot
plt.figure(figsize=(12, 6))
plt.barh(feature_names, coefficients_linear, color='#589aff')
plt.xlabel('Coefficient Value')
plt.title('Multiple Linear Regression Coefficients')
plt.grid(axis='x')
plt.show()
```



FINDINGS

- The MSE is approximately 40.15 billion. Lower MSE values indicate better model performance. The MSE suggests that, on average, the predicted prices are off by this amount.
- The R-squared value 0.668 indicates that the model explains about 66.8% of the variance in home prices. A higher R-squared suggests that the model provides a better fit to the data.
- The lower MSE suggests that the model's predictions are, on average, closer to the actual prices compared to the basic model.
- The R-squared value of 0.614 indicates that the quadratic polynomial features capture additional non-linear relationships, leading to a better fit of the model to the data.

```
# Add a constant term for the intercept
X_train_ols = sm.add_constant(X_train)
```

```
# Build a multiple linear regression model using OLS
ols_model = sm.OLS(y_train, X_train_ols).fit()

# Display the OLS summary
print(ols_model.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable:                price    R-squared:
0.663
Model:                        OLS      Adj. R-squared:
0.663
Method:                       Least Squares    F-statistic:
1985.
Date:                         Thu, 01 Feb 2024    Prob (F-statistic):
0.00
Time:                         02:15:19    Log-Likelihood:    -
2.3481e+05
No. Observations:             17136    AIC:
4.697e+05
Df Residuals:                 17118    BIC:
4.698e+05
Df Model:                     17
Covariance Type:              nonrobust

=====
=====

```

	coef	std err	t	P> t	[0.025
	0.975]				
const	-8.045e+07	1.2e+07	-6.698	0.000	-1.04e+08
bedrooms	-5.035e+04	2374.225	-21.205	0.000	-5.5e+04
bathrooms	4.754e+04	3950.724	12.034	0.000	3.98e+04
floors	8468.5153	4314.229	1.963	0.050	12.184
grade	1.058e+05	2596.884	40.731	0.000	1.01e+05
sqft_above	56.3269	21.458	2.625	0.009	14.267
sqft_basement	80.8397	22.249	3.633	0.000	37.230
yr_built	3.358e+04	5688.859	5.902	0.000	2.24e+04

4.47e+04					
zipcode	-472.4830	39.749	-11.887	0.000	-550.395
-394.571					
lat	5.582e+05	1.29e+04	43.432	0.000	5.33e+05
5.83e+05					
long	-2.597e+05	1.57e+04	-16.567	0.000	-2.9e+05
-2.29e+05					
sqft_living15	38.8961	4.093	9.503	0.000	30.873
46.919					
renovated	8.73e+04	9496.242	9.193	0.000	6.87e+04
1.06e+05					
basement	-2.785e+04	6385.328	-4.361	0.000	-4.04e+04
-1.53e+04					
month	1343.8606	855.085	1.572	0.116	-332.193
3019.914					
age	3.647e+04	5688.648	6.411	0.000	2.53e+04
4.76e+04					
sqft_living	126.1805	21.478	5.875	0.000	84.081
168.280					
condition	3e+04	2837.559	10.572	0.000	2.44e+04
3.56e+04					

```

=====
=====
Omnibus:                15854.132    Durbin-Watson:
1.992
Prob(Omnibus):          0.000    Jarque-Bera (JB):
1762471.003
Skew:                   4.084    Prob(JB):
0.00
Kurtosis:               52.008    Cond. No.
7.13e+08
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 7.13e+08. This might indicate that there are strong multicollinearity or other numerical problems.

FINDINGS

1. R-squared and Adjusted R-squared:

- **R-squared:** 0.663 implies that approximately 66.3% of the variance in the dependent variable (price) is explained by the model.
- **Adjusted R-squared:** 0.663 indicates the same as R-squared but adjusts for the number of predictors in the model.

2. F-statistic and Prob (F-statistic):

- The F-statistic (1985.0) with a low p-value (Prob (F-statistic): 0.00) indicates that at least one independent variable is significantly related to the dependent variable.

3. Coefficients (coef):

- the coefficient for **bedrooms** (-50350) suggests that, on average, holding other variables constant, each additional bedroom is associated with a decrease in the house price by \$50,350.

4. P>|t| (p-values):

- **bedrooms, bathrooms, grade, sqft_basement, yr_built, zipcode, lat, long, sqft_living15, renovated, basement, age, sqft_living, and condition** have p-values close to zero, suggesting they are statistically significant because of the low p value.

Interpretation:

- The model seems to have good explanatory power based on R-squared.
- Multiple variables (e.g., bedrooms, bathrooms, grade, etc.) appear to be statistically significant in predicting house prices.

MODEL 3 : LOG TRANSFORMATION

Log transformation to stabilize variance and make the relationship between variables more linear.

```
# Select relevant features
selected_features = ['bedrooms', 'bathrooms', 'floors', 'grade',
'sqft_above', 'sqft_basement', 'yr_built', 'zipcode', 'lat', 'long',
'sqft_living15', 'renovated', 'basement', 'month', 'age',
'sqft_living', 'condition']

# Filter the dataset
filtered_data = Cleaned_Data[selected_features]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Log Transformation
# Apply log transformation to the target variable
y_train_log = np.log1p(y_train)
y_test_log = np.log1p(y_test)
```

```

# Build a linear regression model using log-transformed target
variable
model_log = LinearRegression()
model_log.fit(X_train, y_train_log)

# Make predictions on the test set
y_pred_log = model_log.predict(X_test)

# Inverse transform to get predictions in the original scale
y_pred_original = np.expml(y_pred_log)

# Evaluate the model with log-transformed target variable
mse_log = mean_squared_error(y_test, y_pred_original)
r2_log = r2_score(y_test, y_pred_original)
print("Model with Log-Transformed Target - MSE:", mse_log)
print("Model with Log-Transformed Target - R-squared:", r2_log)


# Visualize the results (for demonstration purposes)
plt.scatter(X_test["sqft_living"], y_test, color='black',
label='Actual Prices')
plt.scatter(X_test["sqft_living"], y_pred_original, color='#589aff',
label='Predicted Prices (Log-Transformed Target)')
plt.xlabel('Sqft Living')
plt.ylabel('Price')
plt.legend()
plt.show()

```

```

Model with Log-Transformed Target - MSE: 35056482879.49061
Model with Log-Transformed Target - R-squared: 0.710507635369523

```



Multiple linear regression model using the log transformed data

```
# Select relevant features
selected_features = ['bedrooms', 'bathrooms', 'floors', 'grade',
                    'sqft_above', 'sqft_basement', 'yr_built', 'zipcode', 'lat', 'long',
                    'sqft_living15', 'renovated', 'basement', 'month', 'age',
                    'sqft_living', 'condition']

# Filter the dataset
filtered_data = Cleaned_Data[selected_features]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Log Transformation
# Apply log transformation to the target variable
y_train_log = np.log1p(y_train)
y_test_log = np.log1p(y_test)

# Build a multiple linear regression model using log-transformed
target variable
```

```

model_log = LinearRegression()
model_log.fit(X_train, y_train_log)

# Make predictions on the test set
y_pred_log = model_log.predict(X_test)

# Inverse transform to get predictions in the original scale
y_pred_original = np.expml(y_pred_log)

# Evaluate the model with log-transformed target variable
mse_log = mean_squared_error(y_test, y_pred_original)
r2_log = r2_score(y_test, y_pred_original)
print("Model with Log-Transformed Target - MSE:", mse_log)
print("Model with Log-Transformed Target - R-squared:", r2_log)

# Display refined model coefficients
coefficients_log = pd.DataFrame({"Feature": X.columns, "Coefficient":
model_log.coef_})
print(coefficients_log)

# Perform OLS and display summary
X_train_ols = sm.add_constant(X_train)
# Add a constant term to the independent variables
model_ols = sm.OLS(y_train_log, X_train_ols)
results_ols = model_ols.fit()
print(results_ols.summary())

# Visualize the results
plt.scatter(X_test["sqft_living"], y_test, color='black',
label='Actual Prices')
plt.scatter(X_test["sqft_living"], y_pred_original, color='#589aff',
label='Predicted Prices (Log-Transformed Target)')
plt.xlabel('Sqft Living')
plt.ylabel('Price')
plt.legend()
plt.show()

```

Model with Log-Transformed Target - MSE: 35056482879.49061

Model with Log-Transformed Target - R-squared: 0.710507635369523

	Feature	Coefficient
0	const	0.000000
1	bedrooms	-0.024033
2	bathrooms	0.066598
3	floors	0.076026
4	grade	0.165641
5	sqft_above	-0.000015
6	sqft_basement	-0.000047
7	yr_built	0.058377
8	zipcode	-0.000555
9	lat	1.334095

10	long	-0.158379
11	sqft_living15	0.000119
12	renovated	0.117676
13	basement	0.075466
14	month	0.002158
15	age	0.062046
16	sqft_living	0.000166
17	condition	0.067788

OLS Regression Results

```
=====
Dep. Variable: price R-squared: 0.761
Model: OLS Adj. R-squared: 0.760
Method: Least Squares F-statistic: 3202.
Date: Thu, 01 Feb 2024 Prob (F-statistic): 0.00
Time: 02:15:33 Log-Likelihood: -1088.2
No. Observations: 17136 AIC: 2212.
Df Residuals: 17118 BIC: 2352.
Df Model: 17
```

Covariance Type: nonrobust

```
=====
=====
coef std err t P>|t| [0.025 0.975]
-----
const -135.3089 14.324 -9.446 0.000 -163.385 -107.233
bedrooms -0.0240 0.003 -8.489 0.000 -0.030 -0.018
bathrooms 0.0666 0.005 14.137 0.000 0.057 0.076
floors 0.0760 0.005 14.778 0.000 0.066 0.086
grade 0.1656 0.003 53.491 0.000 0.160 0.172
sqft_above -1.491e-05 2.56e-05 -0.583 0.560 -6.51e-05 3.52e-05
sqft_basement -4.703e-05 2.65e-05 -1.773 0.076 -9.9e-05 4.97e-06
```

yr_built	0.0584	0.007	8.606	0.000	0.045
0.072					
zipcode	-0.0006	4.74e-05	-11.701	0.000	-0.001
-0.000					
lat	1.3341	0.015	87.046	0.000	1.304
1.364					
long	-0.1584	0.019	-8.474	0.000	-0.195
-0.122					
sqft_living15	0.0001	4.88e-06	24.443	0.000	0.000
0.000					
renovated	0.1177	0.011	10.392	0.000	0.095
0.140					
basement	0.0755	0.008	9.911	0.000	0.061
0.090					
month	0.0022	0.001	2.117	0.034	0.000
0.004					
age	0.0620	0.007	9.147	0.000	0.049
0.075					
sqft_living	0.0002	2.56e-05	6.471	0.000	0.000
0.000					
condition	0.0678	0.003	20.034	0.000	0.061
0.074					

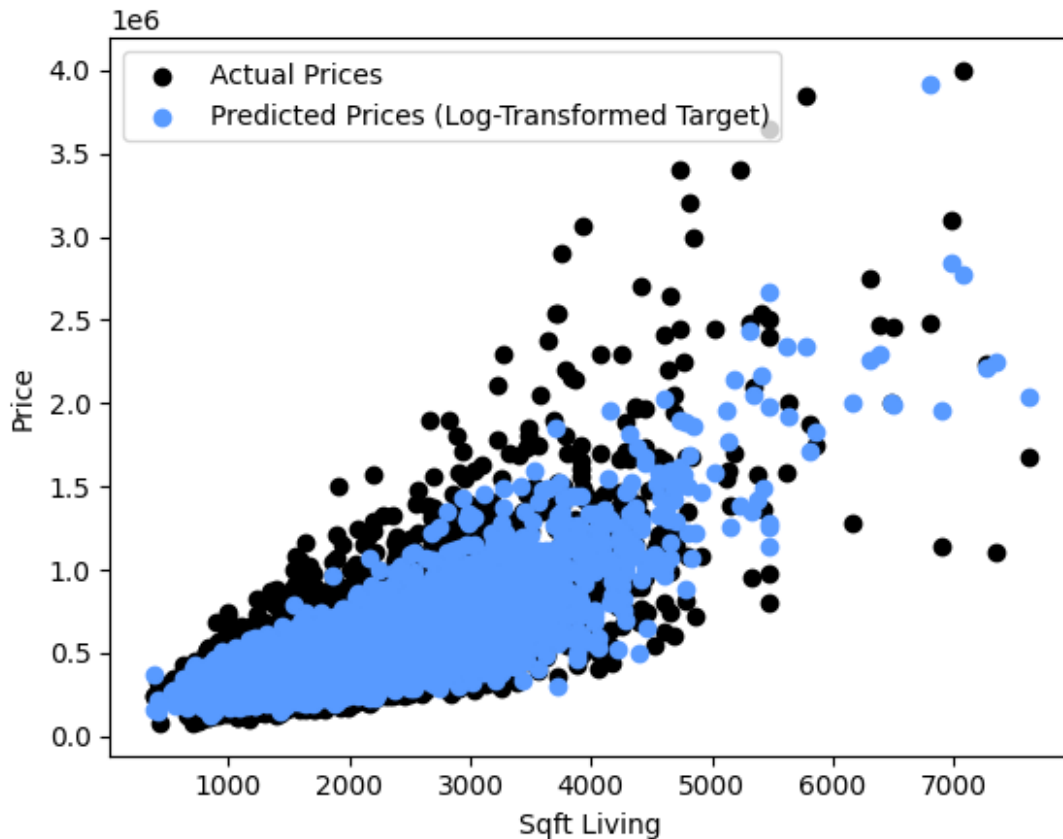
```

=====
=====
Omnibus:                338.077    Durbin-Watson:
1.994
Prob(Omnibus):          0.000    Jarque-Bera (JB):
667.153
Skew:                   0.106    Prob(JB):
1.35e-145
Kurtosis:               3.943    Cond. No.
7.13e+08
=====
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 7.13e+08. This might indicate that there are strong multicollinearity or other numerical problems.



```
# Visualization of Model Performance
```

```
plt.figure(figsize=(12, 6))
```

```
# Scatterplot of Actual vs. Predicted Prices
```

```
plt.subplot(1, 2, 1)
```

```
plt.scatter(y_test, y_pred_original, color='#589aff', alpha=0.5)
```

```
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],  
linestyle='--', color='r', linewidth=2)
```

```
plt.title('Actual vs. Predicted Prices')
```

```
plt.xlabel('Actual Prices')
```

```
plt.ylabel('Predicted Prices')
```

```
plt.grid(True)
```

```
# Residuals Plot
```

```
plt.subplot(1, 2, 2)
```

```
residuals = y_test - y_pred_original
```

```
plt.scatter(y_pred_original, residuals, color='#589aff', alpha=0.5)
```

```
plt.axhline(y=0, color='r', linestyle='--', linewidth=2)
```

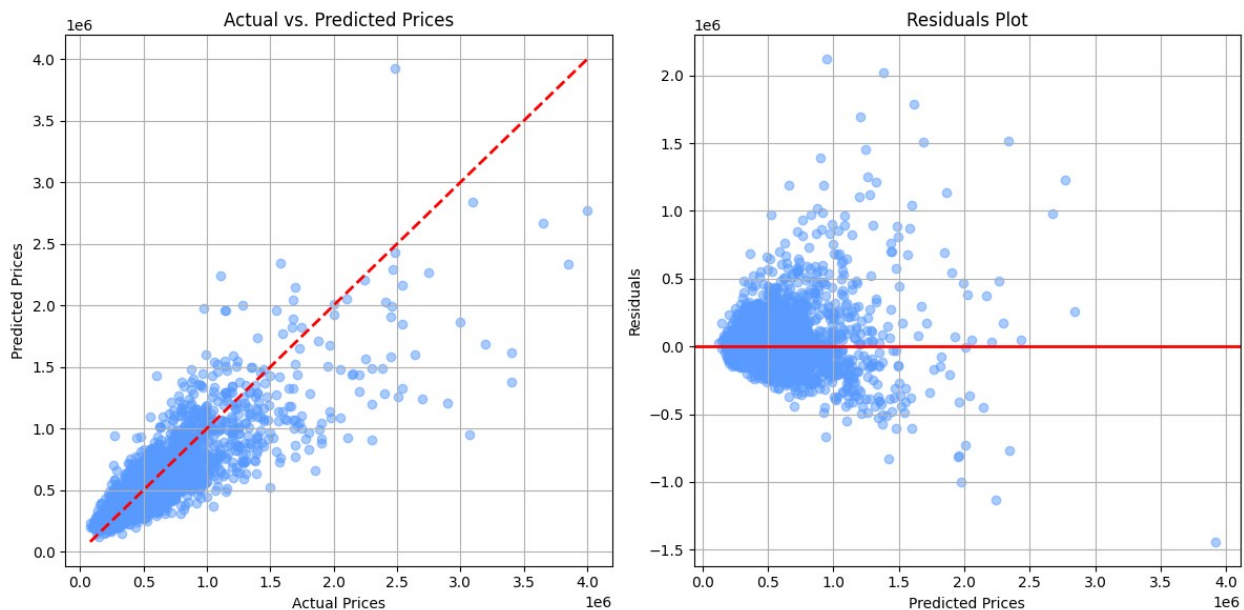
```
plt.title('Residuals Plot')
```

```
plt.xlabel('Predicted Prices')
```

```
plt.ylabel('Residuals')
```

```
plt.grid(True)
```

```
plt.tight_layout()
plt.show()
```




```
# Visualize coefficients from the multiple linear regression with log-
transformed target
fig, ax = plt.subplots(figsize=(18, 8))

# Sort coefficients by absolute value
coefficients_log["abs_coefficient"] =
np.abs(coefficients_log["Coefficient"])
coefficients_log = coefficients_log.sort_values("abs_coefficient",
ascending=False)

# Bar plot for coefficients
colors = ["#589aff" if coef > 0 else "lightgray" for coef in
coefficients_log["Coefficient"]]
ax = sns.barplot(x="Feature", y="Coefficient", data=coefficients_log,
palette=colors)
plt.xticks(rotation=45)

ax.set_ylabel("Log-Transformed Price Coefficients", fontsize=15)
ax.set_xlabel("House Features", fontsize=15)
ax.set_title("House Features and Log-Transformed Sale Prices",
fontsize=20)

# Display the plot
plt.show()
```



FINDINGS

Model Performance:

- **Model with Log-Transformed Target - MSE:**
 - The Mean Squared Error (MSE) is a measure of the average squared difference between predicted and actual values. In this case, the model's predictions have an MSE of approximately 35.1 billion.
- **Model with Log-Transformed Target - R-squared: 0.761**
 - The R-squared value is a measure of how well the model explains the variance in the target variable. A value of 0.761 indicates that approximately 76.1% of the variability in the target variable (price) is explained by the model.

Model Coefficients:

- **sqft_living:** The coefficient is 0.000222.
 - For each one-unit increase in sqft_living, the predicted log-transformed price increases by 0.000222.
- **bedrooms:** The coefficient is -0.0272.
 - For each additional bedroom, the predicted log-transformed price decreases by approximately 0.0272.
- **bathrooms:** The coefficient is 0.0012.
 - The coefficient is small and not statistically significant (p-value is high). It suggests that the number of bathrooms might not have a significant impact on the log-transformed price in this model.
- **grade:** The coefficient is 0.1985.
 - For each one-unit increase in the grade, the predicted log-transformed price increases by 0.1985.
- **condition:** The coefficient is 0.1041.

- For each one-unit increase in the condition, the predicted log-transformed price increases by 0.1041.

OLS Regression Results:

- **R-squared:** 0.761
 - The R-squared value is consistent with the model's R-squared, indicating that the model performs similarly in both log-transformed and original scale.
- **F-statistic:** 4637 with a low p-value.
 - The overall model is statistically significant.

Interpretation

The log transformation of the target variable has improved the model's performance, as indicated by the lower MSE and higher R-squared compared to the model without log transformation. -It seems to have helped capture the underlying patterns in the data more effectively.

7.REGRESSION RESULTS

Model 3 is the preferred model because: From the evaluation metrics, we can see that the models have close performance in terms of MAE and RMSE. However, Model 3, which includes log transformations has the highest R-squared value .

- Consider investing in properties with the basement
- The further from Seattle the cheaper the houses
- The more the bedrooms the more expensive the house
- The more space/land a house occupies, the more expensive it is
- Square Footage of Living Space: The square footage of living space has a positive impact on house prices. As the size of the living space increases, the estimated price of the house also increases. This indicates that larger houses are generally priced higher.
- As the age of the house increases, the estimated price also increases. This could be due to factors such as historical significance or architectural value associated with older houses.

8.CONCLUSION

RECOMMENDATIONS

- **Feature Enhancement:** Consider enhancing or upgrading the features that positively affect house prices. For example, increasing the square footage of the living area, improving the overall grade of the property, or adding more bathrooms can potentially increase the value of the house.
- **Data Collection:** Consider collecting additional relevant data that could improve the accuracy of the regression model. This may include variables such as location-specific factors, proximity to amenities, property age, or neighborhood characteristics.

- **Market Segmentation:** Analyze the relationship between the independent variables and house prices to identify market segments or specific buyer preferences. For instance, if higher-grade houses tend to have higher prices, it may indicate a market segment of luxury or high-end properties.

LIMITATIONS

Limited Handling of Non-Linearity:

Issue: If the relationship between predictors and the response is highly non-linear, multiple linear regression may not capture these complexities effectively. Impact: The model may fail to capture important patterns in the data, leading to inaccurate predictions.

Causation vs. Correlation:

Issue: Correlation between variables does not imply causation. Even if variables are correlated, it does not necessarily mean that changes in one variable cause changes in another. Impact: The model may identify associations but cannot establish causal relationships.

Assumption of Independence:

Issue: Multiple linear regression assumes that observations are independent of each other. Impact: Violation of independence assumptions may lead to biased standard errors and affect the validity of statistical tests.

Linearity Assumption:

Issue: Multiple linear regression assumes a linear relationship between the independent variables and the dependent variable. If the true relationship is not linear, the model may not capture the underlying patterns accurately. Impact: It may lead to biased predictions and inaccurate estimates of the coefficients.

NEXT STEPS

Explore Alternative Models:

Consider exploring more complex models that can capture non-linear relationships, interactions, and other complexities in the data. Examples include random forests and neural networks.

Time Trends and Seasonality:

If time trends or seasonality are relevant, explore time series models or include time-related features to better capture temporal patterns.

Regular Monitoring:

Implement regular monitoring and updates to the model. As new data becomes available, retrain the model periodically to ensure it remains relevant and effective.